

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from statsmodels.tsa.arima.model import ARIMA
5 from sklearn.metrics import mean_squared_error

```

```

1 df = pd.read_csv('/content/Crypto.csv')
2 df.head()

```



	unix	date	symbol	open	high	low	close	Volume XRP	Volume USDT
0	1.640000e+12	12/27/2021	XRP-USDT	0.9200	0.9237	0.9200	0.9226	2384512.0	2.198450e+06
1	1.640000e+12	12/26/2021	XRP-USDT	0.9252	0.9334	0.9052	0.9200	163438501.0	1.499400e+08
2	1.640000e+12	12/25/2021	XRP-USDT	0.9114	0.9350	0.8981	0.9252	250074945.0	2.302303e+08
3	1.640000e+12	12/24/2021	XRP-USDT	0.9941	0.9966	0.8964	0.9115	567234092.0	5.377035e+08
4	1.640000e+12	12/23/2021	XRP-USDT	0.9538	1.0167	0.9372	0.9941	479436230.0	4.729372e+08



Next steps:

[Generate code with df](#)

[View recommended plots](#)
[New interactive sheet](#)

```

1 # Basic Data Exploration
2 print(df.info())

```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1334 entries, 0 to 1333
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   unix        1334 non-null   float64
1   date        1334 non-null   object
2   symbol      1334 non-null   object
3   open        1334 non-null   float64
4   high        1334 non-null   float64
5   low         1334 non-null   float64
6   close       1334 non-null   float64
7   Volume XRP  1334 non-null   float64
8   Volume USDT 1334 non-null   float64
dtypes: float64(7), object(2)
memory usage: 93.9+ KB
None

```

```

1 # Data Cleaning
2 # Convert the 'date' column to datetime
3 df['date'] = pd.to_datetime(df['date'])

```

```

1 # Check for missing values
2 print(df.isnull().sum())

```



```

unix      0
date      0
symbol    0
open      0
high      0
low       0
close     0
Volume XRP 0
Volume USDT 0
dtype: int64

```


```

1 # Plot the closing price over time
2 plt.figure(figsize=(6, 3))
3 plt.plot(df['date'], df['close'])
4 plt.title('Closing Price Over Time')
5 plt.xlabel('Date')
6 plt.ylabel('Closing Price (USDT)')
7 plt.grid(True)
8 plt.show()


```



```
1 # Volume of XRP over time
2 plt.figure(figsize=(6, 3))
3 plt.plot(df['date'], df['Volume XRP'], label='Volume XRP')
4 plt.title('Volume of XRP Over Time')
5 plt.xlabel('Date')
6 plt.ylabel('Volume XRP')
7 plt.legend()
8 plt.grid(True)
9 plt.show()
```




```
1 # Distribution of closing prices
2 plt.figure(figsize=(6, 3))
3 sns.histplot(df['close'], bins=50, kde=True)
4 plt.title('Distribution of Closing Prices')
5 plt.xlabel('Closing Price (USDT)')
6 plt.ylabel('Frequency')
7 plt.show()
```



```

1 # Correlation Heatmap
2 plt.figure(figsize=(5, 4))
3 corr = df[['open', 'high', 'low', 'close', 'Volume XRP', 'Volume USDT']].corr()
4 sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
5 plt.title('Correlation Matrix')
6 plt.show()

```



```

1 # Time Series Modeling with ARIMA
2 # Split the data into train and test sets
3 train_size = int(len(df) * 0.8)
4 train, test = df['close'][:train_size], df['close'][train_size:]

```

```

1 # Fit the ARIMA model
2 model = ARIMA(train, order=(5, 1, 0))
3 model_fit = model.fit()

```

```

1 # Print the model summary
2 print(model_fit.summary())

```

```

=====
SARIMAX Results
=====
Dep. Variable:      close      No. Observations:      1067
Model:              ARIMA(5, 1, 0)  Log Likelihood      1782.988
Date:               Thu, 25 Jul 2024  AIC              -3553.977
Time:               17:10:04      BIC              -3524.147
Sample:             0            HQIC              -3542.674
Covariance Type:    opg
=====
              coef      std err      z      P>|z|      [0.025      0.975]
-----
ar.L1      -0.0644      0.018      -3.565      0.000      -0.100      -0.029
ar.L2       0.0021      0.016       0.131      0.896      -0.029      0.033
ar.L3       0.0165      0.012       1.354      0.176      -0.007      0.040
ar.L4       0.0618      0.015       4.235      0.000      0.033      0.090
ar.L5      -0.0339      0.014      -2.494      0.013      -0.061      -0.007
sigma2       0.0021      2.62e-05      78.675      0.000      0.002      0.002
=====
Ljung-Box (L1) (Q):      0.01  Jarque-Bera (JB):      36131.30
Prob(Q):                0.93  Prob(JB):              0.00
Heteroskedasticity (H):  0.03  Skew:              0.05
Prob(H) (two-sided):    0.00  Kurtosis:          31.52
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```


```

1 # Make predictions
2 start = len(train)
3 end = len(train) + len(test) - 1
4 predictions = model_fit.predict(start=start, end=end, typ='levels')

```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/representation.py:374: FutureWarning: Unknown keyword arguments: dict
warnings.warn(msg, FutureWarning)
```

```
1 # Plot the predictions against the actual values
2 plt.figure(figsize=(12, 6))
3 plt.plot(df.index, df['close'], label='Actual')
4 plt.plot(df.index[start:end+1], predictions, color='red', label='Predicted')
5 plt.title('ARIMA Model - Actual vs Predicted')
6 plt.xlabel('Date')
7 plt.ylabel('Closing Price (USDT)')
8 plt.legend()
9 plt.grid(True)
10 plt.show()
```



```
1 # Calculate and print the mean squared error
2 mse = mean_squared_error(test, predictions)
3 print(f'Mean Squared Error: {mse}')
```

```
Mean Squared Error: 0.03538190888201755
```